

1. To implement the `imageGradientX()` function I used `numpy.ndenumerate()` function to iterate over the image. This function returns an iterator with pairs of array indices and values. I used the indices to capture the value of the next pixel in the x direction, which I then used to calculate the change. I set the pixels along the edges to 0 indicating that there is no change because there are no more pixels to compare. After calculating the gradient I used `numpy.absolute()` function to convert that to the magnitude of the gradient.

```
image (numpy.ndarray): A grayscale image represented in a numpy array.

Returns:
    output (numpy.ndarray): The image gradient in the X direction.
"""
# WRITE YOUR CODE HERE.
y = 0
x = 1

output = np.zeros(image.shape)
image = image.astype np.float64)

for index, pixel in np.ndenumerate image):
    # Edge case
    if index[x] == image.shape[x] - 1:
        output[index[y], index[x]] = 0
    else:
        output[index[y], index[x]] = image[index[y], index[x] + 1] - pixel

return np.absolute output)
# END OF FUNCTION.

def imageGradientY image):
    AG Get Help    AO WriteOut    AR Read File    AY Prev Page    AK Cut T
    AX Exit        AJ Justify     AW Where Is    AW Next Page    AU UnCut
```

2. The `imageGradientY()` function was almost exactly the same except I used the indices to capture the value of the next pixel in the y direction, which I then used to calculate the change. I found I needed to convert the image array to `float64` when my initial attempts resulted in noise. This was due to overflow errors as the input array was `uint8`.

```
Args:
    image (numpy.ndarray): A grayscale image represented in a numpy array.

Returns:
    output (numpy.ndarray): The image gradient in the Y direction.
"""
# WRITE YOUR CODE HERE.
y = 0
x = 1

output = np.zeros(image.shape)
image = image.astype np.float64)

for index, pixel in np.ndenumerate image):
    # Edge case
    if index[y] == image.shape[y] - 1:
        output[index[y], index[x]] = 0
    else:
        output[index[y], index[x]] = image[index[y] + 1, index[x]] - pixel

return np.absolute output)
# END OF FUNCTION.
```

3. To implement the computeGradient() function I again used the numpy.ndenumerate() function to iterate through the image as well as the kernel. I first checked the bounds of the image to ensure the cross-correlation would take place in the bounds of the image. I then looped through the kernel using the image and kernel indices to capture the proper pixels to perform the cross-correlation on.

```
Returns:
    output (numpy.ndarray): The computed gradient for the input image.
"""
# WRITE YOUR CODE HERE.
output = np.zeros(image.shape)
image = image.astype(np.float64)

y = 0
x = 1

for imageIndex, imagePixel in np.ndenumerate(image):
    # offset the kernel to prevent out of bounds exception
    if imageIndex[x] > 0 and imageIndex[y] > 0 \
        and imageIndex[x] < image.shape[x] - 1 \
        and imageIndex[y] < image.shape[y] - 1:

        sum = 0
        for kernelIndex, kernelValue in np.ndenumerate(kernel):
            u = kernelIndex[x] - 1
            v = kernelIndex[y] - 1

            sum += image[imageIndex[y] + v, imageIndex[x] + u] * kernelValue

        output[imageIndex[y], imageIndex[x]] = sum

return output
# END OF FUNCTION.
```

AG Get Help AO WriteOut AR Read File AY Prev Page AK Cut Text
AX Exit AJ Justify AW Where Is AN Next Page AU UnCut Text

4. I started with the following image while experimenting with edge detection:



John Radice
GTID 903095161
Assignment 5

After experimenting with several kernels, left/right average derivative, Prewitt, Sobel, Gaussian derivative, etc. I found the Sobel kernel provided the most accurate depiction of the edges. After taking the X and Y gradients I converted them to black and white. I used a threshold of 96 when converting the gradient in the X direction and used a threshold of 128 when converting the gradient in the Y direction. To create the final image shown below I averaged the X and Y gradients after applying the threshold.



For reference I have included the cv2.Canny() result of my original image:

